

PROJECT DOCUMENTATION

Counselor NIDS

A network intrusion detection system in which machine-learning detectors that advise each other and the Snort rule engine watch the same live traffic, and in which Snort can guide the detectors on a network they were never trained on.

Final-year thesis project · October 2026

Source code: github.com/hash-cmd/counselor-nids

Contents

1. Summary
2. The problem and the thesis objectives
3. System design
4. The AI detectors and the counselors network
5. Snort alongside the AI
6. What happens to each connection
7. Training the detectors
8. Results
9. The thesis extension: Snort-guided adaptation
10. The dashboard and the API
11. Deployment, operations and security
12. Limitations
13. Project layout and how to run it
14. References

1. Summary

Counselor NIDS watches the traffic on a network interface and raises an alarm when a connection looks like an attack. Two independent checks judge every connection:

- **Three machine-learning detectors**, each a specialist in one family of attacks: floods (DoS), break-ins (password guessing and website attacks) and botnets. When a detector is unsure, it asks the others for advice. This is the *counselors network* of Quincozes et al. (2019).
- **Snort 3**, a widely used signature engine that recognises known attack patterns inside the traffic.

Every Snort alert is linked to the connection it belongs to and compared with the AI's verdict on that connection, so each incident is reported as *Confirmed* (both agree), *Disputed* (Snort only) or *AI only*. A web dashboard shows the incidents, who is attacking, how the detectors reached their decisions, how many false alarms the system raises on the network over days, and whether every part of the system is running.

On real attack traffic it never trained on, the system catches 99–100% of most attack types, including floods, password guessing and botnet traffic, with almost no false alarms (Section 8). The thesis extension (Section 9) adds Snort as a counselor, self-learning from agreement between Snort and the AI, and adaptive trust, so the detectors can adapt to a network they were never trained on without anyone labelling its traffic.

2. The problem and the thesis objectives

Machine-learning intrusion detectors score almost perfectly on the dataset they were trained on, but lose most of that accuracy on a different network. Trained on CIC-IDS2017 and tested on CSE-CIC-IDS2018, published classifiers fall to about chance accuracy (Cantone et al., 2024), or a macro-F1 near 0.12. Every real deployment is a "different network", and labelling its traffic by hand is too expensive, so this is the main reason AI-based detection is not trusted in practice.

GENERAL OBJECTIVE

To design, implement and evaluate a counselor-based network intrusion detection system in which a signature-based IDS (Snort) guides machine-learning detectors to adapt to a network they were never trained on, without labelled traffic from that network.

Specific objective	Challenge it addresses
1. Snort as a counselor	In the counselors network, detectors only ask other AI detectors, who can share the same blind spots on a new network. Snort's signatures behave the same on any network, so they give independent advice.
2. Learning from agreement	There are no labels on the new network. Connections where Snort and the AI independently agree serve as labels for free, and a safety gate blocks retraining that hurts or makes the detectors forget what they knew.
3. Adaptive trust	Each counselor's trust is measured once, in the lab. It is re-estimated on the new network from how often the others agree with it: for every Snort rule, and (tested as an ablation) for the AI detectors themselves.
4. Cross-network evaluation	The mechanisms are evaluated by training on one lab (CSE-CIC-IDS2018) and deploying on another (CICIDS2017), in both directions, against each baseline.
5. Live deployment	Research detectors rarely run on live traffic. This one does, and the practical obstacles found along the way are documented.

3. System design




Figure 1. System design: two lanes judge the same packets; Redis carries the messages between services; the magenta row is the thesis addition.

Service	Command	Role
Traffic capture (Extractor)	<code>nids extract -live IFACE</code>	Captures packets and turns them into connection records (flows) with the Python <code>cicflowmeter</code> . Flows that carry no data go to Snort only. Survives the interface going down and resumes on its own.
Observer	<code>nids observe</code>	Routes each batch of flows to every detector subscribed to that traffic source.
Detector (×3)	<code>nids detect MODEL</code>	Classifies flows, asks the other detectors for advice and answers their requests, and publishes alerts.
Snort bridge	<code>nids snort --interface IFACE</code>	Runs Snort on the same interface, links each alert to its connection and asks the detectors for their verdict on it.
Alert journal	<code>nids journal</code>	Keeps every alert on disk (one file per day) to measure false alarms over days.
Health	<code>nids health</code>	Every service sends a heartbeat; this lists them and exits with an error if any part has stopped.
API and dashboard	Django (Daphne), Next.js	REST and WebSocket API and the web dashboard.

TECHNICAL NOTE

All services communicate through Redis streams: `nids:unknown` (new flows), `nids:inbox:<detector>`, advice request and reply lists, `nids:alerts` and `nids:snort:alerts`, a flow index (sorted sets keyed by connection and by host pair), and a `nids:health` hash of heartbeats. Because each part is a separate process, any one can crash and restart without stopping the others.

4. The AI detectors and the counselors network

How one detector decides (from the paper)

- Training (Algorithm 1).** Five classifiers are trained: naive Bayes, three kinds of decision tree (J48-like, REPTree-like, RandomTree-like) and a random forest. The evaluation data is grouped into clusters with K-Means. For every cluster, the most accurate classifier, and any within 0.1% of it, is kept.
- Detection (Algorithm 2).** A new connection is placed in its nearest cluster, and that cluster's selected classifiers vote. If they agree, that is the verdict. If they disagree, there is a *conflict*.
- Counselors.** On a conflict, the detector asks the other detectors (its counselors) for their verdict on the same connection. It takes the advice of the most confident counselor, if that counselor's accuracy is at least 90%.

Extensions made in this project

- **Cross-check.** A detector cannot raise a conflict about an attack family it never learned; it is just confidently wrong. So every "normal" verdict is also checked with the counselors, and a confident "attack" advice overrides it.
- **Unresolved conflicts count as normal.** When the classifiers disagree and no counselor can help, the paper falls back on the best classifier's guess. On live traffic those guesses caused most false alarms while accounting for almost no real detections, so they are treated as normal.
- **Connection-only flows skip the AI.** Port scans and failed connections exchange no data. The detectors never learned them and only produced guesses (one scan caused ~50,000 false alarms), so they are left to Snort.

Detector	Model file	Learned
Flood detector	<code>models/live/live_dos.joblib</code>	DoS GoldenEye, Hulk, Slowloris, SlowHTTPTest (and DDoS once the 2017 data is added)
Break-in detector	<code>models/live/live_access.joblib</code>	FTP and SSH password guessing, website attacks
Botnet detector	<code>models/live/live_bot.joblib</code>	Ares-infected hosts calling their command-and-control server

5. Snort alongside the AI

Snort 3 runs on the same interface with the project's 42 rules (`snort/rules/nids.rules`) and the ~4,000 Snort 3 community rules. Each alert is linked to its connection: the same two addresses and ports, the same protocol, and a time inside the connection's start-to-end window. A port-scan alert stands for many probes, so it is linked to every connection between the two hosts within ± 30 seconds. The detectors are then asked for their verdict on that connection, using the same advice request they use with each other.

Agreement	Meaning
Confirmed	Snort and the AI both say attack: very likely real.
Disputed	Snort says attack, the AI says normal: a possible Snort false alarm, or an attack the AI never learned (for example SQL injection, which looks like ordinary web traffic to the AI).
AI only	The AI flagged a connection Snort had no alert for: typically floods and botnet traffic.

Rules and tuning

- **Website attacks:** SQL injection (UNION SELECT; a quote followed by an AND/OR comparison; ORDER BY probing; `information_schema`), cross-site scripting (`<script` , event handlers, `javascript:`) in the URL or a form, and website login brute force (20 password posts from one source in 60 s).

- **FTP and SSH brute force** over a wide 120 s window, so slow guessing paced to slip under a short-window threshold is still caught. **SYN and HTTP floods** are covered by rate thresholds, **port scans** by Snort's scan inspector.
- **Checksums are not checked.** Network cards fill checksums in themselves (offloading), so a machine's own outgoing packets look broken to Snort, which used to drop them. Until this was fixed, Snort raised no website-attack alert at all on the 2018 website-attack capture.
- **Quieted rules:** protocol-anomaly alerts are off, and informational community rules (pings, ICMP errors, 1-minute DNS TTLs, UPnP discovery) are suppressed. They fired ~3,800 times on a normal workstation.

6. What happens to each connection




Figure 2. The decision path of one connection through the AI (left) and Snort (right), and the thesis learning loop (bottom).

7. Training the detectors

KEY FINDING

Detectors trained on the published dataset CSVs caught **0%** of password-guessing attacks measured the live way. The CSVs were made with the Java CICFlowMeter; live capture uses the Python cicflowmeter, and the two differ in 71 of 154 feature medians (packet lengths with or without headers, flag counting, where flows are split). So the detectors are trained on the datasets' **raw packet recordings**, measured by the same flow meter as live capture.

1. **Fetch:** single hosts' captures from the CSE-CIC-IDS2018 archive (about 1.3 GB, using HTTP range requests, out of 40–60 GB archives).
2. **Label:** the attack windows are found in the packets themselves, and each flow is labelled by its attacker-victim pair. The dataset's CSV timestamps can't be used: they're a 12-hour clock without AM/PM.
3. **Measure:** the Python flow meter, reading recordings by their own clock (Section 12 explains the bug this fixed). Snort also runs on every slice, and the rules that fired are kept with each flow.
4. **Split by time:** within every attack type and capture, the earliest 80% of flows train and the latest 20% test, so test flows come from later in the attack than anything learned.
5. **Promotion gate:** the installed detectors are tested on the same test flows. New ones replace them only if false alarms don't rise and no attack type drops by more than 2 points. Model files are pickles, which run code when loaded, so the installed set is listed with SHA-256 checksums that live mode verifies before loading.

8. Results

Test split: the latest 20% of every attack (CSE-CIC-IDS2018)

Traffic	Test flows	Detected by the system
DoS SlowHTTPTest	21,110	100.00%
DoS Hulk	5,440	99.94%
FTP brute force	4,416	100.00%
Botnet	3,432	99.83%
DoS GoldenEye	3,405	99.88%
DoS Slowloris	2,166	84.53%
SSH brute force	1,015	99.01%
Website attacks	42	92.86%
Normal traffic (false alarms)	5,192	0.04% (2 flows)
All attacks		99.12%

Whole system on a held-out recording

`real-attacks-2018.pcap` holds the final minutes of every attack, which the detectors never trained on, plus normal office traffic. It runs through the flow meter, the detectors, Snort with the project's rules, and the alert linking.

Traffic	Flows	AI	Snort	Either
DoS GoldenEye	1,822	99.1%	4.9%	99.1%
DoS Hulk	2,861	99.9%	48.2%	99.9%
DoS SlowHTTPTest	4,450	100.0%	99.8%	100.0%
DoS Slowloris	760	34.7%	0.0%	34.7%
FTP brute force	5,040	100.0%	99.8%	100.0%
SSH brute force	359	98.3%	88.9%	98.9%
Website attacks	4	4 of 4	4 of 4	4 of 4
Botnet	540	99.6%	0.0%	99.6%
Normal (false alarms)	1,570	0.0%	0.3%	0.3%

- **The two checks cover each other.** The AI catches floods and the botnet that Snort's rules miss. Snort covers website attacks: on the whole 2018 website-attack capture it alerts on 188 of the attacker's 207 connections (90.8%) and on nothing else.

- **A botnet host it never saw:** 99.8% of 15,238 botnet flows caught, with false alarms on 2 of 4,307 normal flows.
- **Slowloris is the weak spot,** and the rate depends on the attack's phase: 85% across the attack, 60% over its complete last five minutes, 35% in the last three as the attack winds down. Snort's Challenge-ACK community rule lifts the last three minutes to 69.5%.

9. The thesis extension: Snort-guided adaptation

9.1 Snort as a counselor

The counselors network now accepts *advisors*: counselors that give advice but classify nothing themselves. Snort is one. When a detector's classifiers disagree on a connection (a conflict), Snort's verdict on the same connection is advice like any other, worth the trust of the most trusted rule that fired. Snort only ever says "attack", and only on connections it alerted on: silence is not evidence of normal traffic, since Snort cannot see floods or botnets.

- **Conflicts only.** Snort does not take part in cross-checking confident "normal" verdicts. In development runs that turned the system into "AI or Snort", inheriting every Snort false alarm (27% of normal connections, mostly Windows policy rules).
- **A trusted "attack" wins.** On a conflict, the most trusted "attack" advice beats any "normal" advice. Otherwise the flood detector, which knows nothing about website attacks, could confidently overrule Snort on one. This is the same "limited vision" reasoning as the cross-check.

9.2 Learning from agreement

While the detectors watch a new network, each chunk of traffic produces *agreement labels*:

- **attack:** Snort alerted and the AI says attack *without Snort's advice* (so Snort cannot confirm itself);
- **normal:** Snort is silent and every detector is unanimously sure the connection is normal.

Everything else stays unlabelled. Each detector stays a specialist: it learns "attack" labels only for connections it was involved in (it said attack, or was unsure), and every detector learns the "normal" labels. A retrained detector replaces the old one only if:

1. it is at least as accurate on a held-back fifth of the agreement labels; and
2. it has not **forgotten**: on the original lab's labelled test data, no attack type it knew drops by more than 2 points and false alarms don't rise. Without this check, retraining on mostly flood labels made the break-in detector lose website attacks.

No label from the new network is used.

9.3 Adaptive trust

- **Snort rules (in the proposed system):** each rule's trust is a Beta estimate (prior 0.9, worth 20 observations), raised when the AI independently agrees and lowered when every detector is sure the connection is normal. In development runs, the Windows RDP policy rule fell to 0.02 trust after the AI contradicted it 786 times, while the attack rules rose to 0.996–0.9998.
- **AI detectors (ablation):** each cluster's accuracy, its advice confidence and optionally its choice of classifiers, re-estimated from agreement labels with the lab value as a prior worth 200 samples. This turned out harmful. Agreement labels are the easy cases every classifier

gets right, so the estimates are biased upwards: re-selecting classifiers let weaker ones vote (more false alarms, fewer website attacks caught), and inflated confidence let "normal" advice win cross-checks it should lose. Agreement can safely re-weight Snort's rules, but not the AI's own classifiers.

9.4 The experiment

`scripts/live_detectors/adapt.py` trains the detectors on one lab and streams the other lab's traffic through the system in time order, in chunks. A fixed test set (the target's latest 20%) is scored before adaptation and after every chunk. It compares:

Configuration	Uses
Snort alone · AI or Snort	Baselines
<code>ai</code>	The counselors network as deployed today
<code>+snort</code>	+ Snort as a counselor
<code>+rules</code>	+ Snort counselor and adaptive rule trust
<code>+learn</code>	+ Snort counselor and learning from agreement
<code>full</code>	The proposed system: Snort counselor + rule trust + learning
<code>+trust</code> , <code>+trust-reselect</code> , <code>full+trust</code>	Ablations with adaptive detector trust

It reports detection rate, false-alarm rate, per-attack detection and a balanced score at every step, each rate with a 95% Wilson confidence interval; McNemar's test for every comparison with the deployed system and the baselines (is the difference real or chance?); how many agreement labels each chunk produced and how many were right; and the final trust of every Snort rule.

9.5 First results (development run, single lab)

A development run on CSE-CIC-IDS2018 alone (the detectors adapt to and are tested on its held-out part) checks the mechanics. It is not the cross-network result.

Configuration	Slowloris	Website attacks	All attacks	False alarms	McNemar vs ai
<code>ai</code> (deployed)	69.1%	71.4%	98.23%	0.08%	—
<code>+snort</code>	93.6%	95.2%	99.57%	0.08%	275 vs 0, $p < 10^{-60}$
<code>+learn</code>	98.0%	71.4%	99.83%	0.08%	332 vs 4, $p < 10^{-70}$
<code>+trust-reselect</code> (ablation)	94.3%	52.4%	99.57%	0.42%	279 vs 14

The agreement labels were right 100% of the time for attacks (15,941 labels) and 99.9% for normal traffic. Snort as a counselor and learning from agreement both raise detection significantly with no extra false alarms.

STATUS

All three mechanisms and the experiment harness are implemented. The cross-network results need the CICIDS2017 packet captures (about 50 GB, being downloaded); the builder for that data also runs Snort on every day automatically. Results will be added when the runs complete.

9.6 In live mode

The mechanisms also run in the live system. Whenever Snort runs, the detectors start with `--snort-counselor`; `NIDS_SNORT_COUNSELOR=0` turns it off.

- **Linking:** the Snort bridge links each alert to its connection as soon as the connection is known.
- **Waiting for Snort:** a detector that is unsure waits until Snort has seen that connection to the end (at most 10 seconds), then takes its advice.
- **Rule trust:** adapts from Confirmed and Disputed verdicts, and is kept in `logs/snort-trust.json` between runs.
- **Fair disagreements:** a disagreement counts against a rule only if the AI flagged nothing from the alert's source. Without this, the Challenge-ACK rule, which catches the Slowloris connections the AI misses, lost all its trust (0.095); with it, the rule keeps 0.97, while the noisy Windows policy rules still drop to 0.09 and 0.36.

On the held-out recording, the live services with Snort as a counselor caught **47% of Slowloris, against 35% without**, with no false alarms (community rules on, live adaptive trust). Offline, with fixed trust, the same recording reaches 69.5%. In full-speed replays, how many doubts Snort settles varies between runs, because the services race each other; live traffic arrives in real time.

10. The dashboard and the API

Page	What it shows
Overview	What is being watched, and the health of every service; headline numbers with a plain-language verdict; top-risk attackers; AI and Snort detections over time; false alarms on your network, with attack-test marking; the latest incidents
Alerts	Every flagged connection once, with a severity and a plain explanation. An Attackers view groups alerts by source into ranked attack stories, with full-history search and a Print/PDF report
Detectors	Each detector's counters and how its decisions were made: on its own, by advice, double-checked or a best guess; Snort as a counselor, with how many doubts it settled and its live trust in every rule
Test with recordings	Replay a recording through the AI and Snort, apart from live monitoring
Results	The detectors' test results per attack type with 95% intervals; the whole system on a held-out recording, with and without Snort as a counselor; the unseen botnet machine; and the thesis experiments (score over time, every configuration with intervals and McNemar, agreement-label accuracy, rule trust)

- Written for non-experts ("AI" and "rule checker" rather than "ML" and "Snort").
- A notification bell for new high-severity alerts.
- **Block this IP** writes out a ready-to-paste firewall command; it never runs anything itself.
- Offline IP reputation from public blocklists, looked up locally, so no address leaves the machine.

The API offers login with httpOnly cookies (plus JWT tokens for scripts), detector counters and service health, alerts and incidents, Snort comparisons, the false-alarm journal and attack-test marking, results, IP reputation, replay control, and a WebSocket live feed that sends stats, alerts and health every second.

11. Deployment, operations and security

- **One command:** `./start.sh setup`, then `./start.sh live wlan0`, or `./start.sh --prod live wlan0` for production servers.
- **As services:** `sudo deploy/install-systemd.sh wlan0` installs one systemd unit per part. Each restarts on failure, and the detectors restart with capture so two runs never mix. Nothing runs as root: capture and Snort get only `CAP_NET_RAW` and `CAP_NET_ADMIN`.
- **Health:** every service sends a heartbeat. The dashboard shows one chip per part, and `nids health` runs every minute from a timer.
- **Logs:** the system journal under systemd. Under `start.sh`, logs over 20 MB are rotated. The alert journal keeps one year.
- **Security:** httpOnly SameSite=Strict session cookies; a custom header on every state-changing request as a CSRF guard; at most 10 login attempts a minute per client. In

production the API refuses to start without a real secret key, and it sends security headers. Model files are checksum-verified before loading.

12. Limitations

- The detectors learned from one lab with one attacker per attack type. This is exactly the problem the thesis extension tackles; until it is evaluated, check false alarms on your own network with the journal.
- The botnet detector knows one botnet family (Ares). There is no live DDoS or infiltration detector yet; DDoS comes with the 2017 data.
- Slowloris detection varies with the attack's phase (35–85%).
- Website attacks rely mainly on Snort: the AI had only about 200 website-attack flows to learn from.
- Snort-guided adaptation can only use Snort's help for attacks Snort can see. For floods and botnets it depends on the AI detectors agreeing with each other.

13. Project layout and how to run it

```
nids/
├─ start.sh          set up and run everything
├─ deploy/           systemd units and their installer
├─ docs/             deployment, live testing, datasets, diagrams, these PDFs
├─ models/live/      the three detectors, manifest.json, SHA256SUMS
├─ snort/            Snort 3 configuration and the project's rules
├─ src/nids/         capture, detector, counselor (incl. Snort counselor), services,
cli
├─ scripts/live_detectors/ fetch, build, train, evaluate, adapt (the thesis experiment)
├─ backend/          Django API: REST, WebSocket live feed, replays, login
└─ frontend/         Next.js dashboard
```

```
./start.sh setup          # once
./start.sh live wlan0     # watch the network, dashboard on http://localhost:3000
.venv/bin/nids health     # are all parts running?
.venv/bin/nids journal-report # false alarms over days
python scripts/live_detectors/adapt.py --source 2018 --target 2017 # thesis experiment
```

14. References

- V. E. Quincozes et al., "A Counselors-Based Intrusion Detection Architecture", 2019; and "Counselors network for intrusion detection", International Journal of Network Management, 2020 (doi:10.1002/nem.2111).
- I. Sharafaldin, A. H. Lashkari, A. A. Ghorbani, "Toward Generating a New Intrusion Detection Dataset and Intrusion Traffic Characterization", ICISP 2018 (CICIDS2017); CSE-CIC-IDS2018, Canadian Institute for Cybersecurity.

- M. Cantone et al., "On the Cross-Dataset Generalization of Machine Learning for Network Intrusion Detection", 2024 (arXiv:2402.10974).
- G. Engelen, V. Rimmer, W. Joosen, "Troubleshooting an Intrusion Detection Dataset: the CICIDS2017 Case Study", WTMC 2021; M. Lanvin et al., "Errors in the CICIDS2017 Dataset and the Significant Differences in Detection Performances It Makes", 2023.
- Snort 3 and the Snort 3 community rules, Cisco Talos (snort.org).