Counselor NIDS

How Each Methodology Step Was Done, and How the Code Does It

For a colleague reading the project: each stage explains what we did and the calculation, the file and code that carries it out line by line, and a worked example with real numbers. Every stage is tagged with the objective and research question it serves, so this document also shows how the methodology answers the problem statement, the objectives and the research questions.

What the project solves (referenced by the tag on every stage below)

Problem. Existing systems combine rules and machine learning too simply and can teach themselves to ignore new attacks, so previously unseen attacks are missed while false positives stay high and alerts are unexplained.

Objectives. O1 design the architecture · O2 implement it on corrected CIC-IDS2017/2018 with one flow extractor · O3 evaluate unseen-attack detection (leave-one-family-out, cross-dataset) · O4 measure verification's effect on false positives · O5 assess explanation quality.

Research questions. RQ1 detect unseen attacks vs baselines · RQ2 does independent verification cut false positives without losing unseen attacks · RQ3 do the novelty score and "AI-only" find unseen attacks · RQ4 are the explanations faithful and plausible.

1. Data collection and flow extraction

Addresses → Objective O2 — same flow extractor for training and live operation.

Method. A capture is a raw stream of packets, but a model needs a fixed-size description of each connection, so we reconstruct connections and summarise them. Packets are grouped by their 5-tuple (source/destination address, source/destination port, protocol) into bidirectional flows, and each flow is turned into ~80 numbers: flow duration = (last packet time – first packet time); total forward packets = count of packets from source to destination; inter-arrival features summarise the gaps between successive packets; TCP-flag counts tally how many packets carried SYN, ACK, FIN, etc. The same extractor is used for the datasets and for live capture.

In code → src/nids/capture/flowmeter.py (Python cicflowmeter) · src/nids/datasets/flow_features.py

```
features = list(CIC2018_TO_2017.values())    # the ~77 shared flow features
# the flow meter groups packets by 5-tuple and emits one row of features per flow
```

The code, step by step:

- `CIC2018_TO_2017` maps the 2018 feature names onto the 2017 names, giving one shared feature vocabulary so both datasets describe a flow the same way.
- The flow meter (a patched cicflowmeter) does the grouping and the arithmetic; its output is one DataFrame row per connection, with a label column.

Worked example. A connection sends 20 packets between 10:00:00.000 and 10:00:04.000. The extractor records flow duration = 4.000 s, total packets = 20, and (if 12 went one way) fwd packets = 12 — three of its ~77 numbers.

2. Cleaning and standardising

Addresses → Objective O2 · underpins all RQs — removes duplicates and prevents leakage.

Method. Rate features (bytes ÷ duration) can be infinite when a flow lasts almost no time — those rows are dropped. Exact-duplicate rows are removed so one flow cannot sit in both training and test. Then each feature is put on a common scale: a signed logarithm first compresses the range, and a standardiser maps $x \to (x-\mu)/\sigma$. The scaler's μ and σ are learned on training rows only.

In code → scripts/live_detectors/data_prep.py : prepare() · src/nids/detector/preprocessing.py

```
flows = flows.replace([np.inf,-np.inf], np.nan).dropna(subset=features)
flows = flows.drop_duplicates(subset=features)          # remove exact duplicates
flows['is_attack'] = flows['label'] != 'Benign'         # binary target
# build_preprocessor(): signed_log1p -> StandardScaler -> VarianceThreshold (fit on train)
```

The code, step by step:

- Line 1 turns infinities into NaN and drops those rows — the undefined rates.
- Line 2 compares the whole 77-feature vector, so only genuine repeats are removed (about one flow in six in our data).
- Line 3 collapses the many attack labels into a single attack/benign target.

- The preprocessor log-compresses, then standardises, then drops constant columns; fitting it inside `DetectorModel.fit` on training keeps the test set's statistics out of training.

Worked example. A flow with 5,000 bytes: signed-log gives ln(1+5000) ≈ 8.52. If, across training flows, that logged feature has mean 6.0 and std 1.2, the standardised value is (8.52−6.0)/1.2 ≈ +2.1 — this flow is 2.1 standard deviations above average on that feature.

3. Three-way, time-ordered split (60 / 20 / 20)

Addresses → Underpins all RQs — the leakage-free split that answers the problem statement's "weak evidence" gap.

Method. Within each (attack type, capture day) we sort flows by start time and cut at 60% and 80%: the earliest 60% is training, the next 20% validation, the last 20% test. Splitting by time means the model never trains on traffic recorded after what it is tested on; splitting within each label means every attack family appears in all three sets. The test set is used once.

In code → scripts/live_detectors/data_prep.py : three_way_split()

```
flows = flows.sort_values('flow_start', kind='stable')
for _, g in flows.groupby(['label','capture']):
    a, b = int(len(g)*0.6), int(len(g)*0.8)
    train.append(g.iloc[:a]); val.append(g.iloc[a:b]); test.append(g.iloc[b:])
```

The code, step by step:

- The stable sort keeps flows with the same timestamp in a fixed order, so the split is reproducible.
- Grouping by (label, capture) and cutting each group guarantees the 60/20/20 proportions per family, not just overall.
- `iloc[:a] / iloc[a:b] / iloc[b:]` take the earliest, middle and latest slices.

Worked example. A day has 10,000 DDoS flows. a = 6,000 and b = 8,000, so the earliest 6,000 train the model, the next 2,000 tune its thresholds, and the latest 2,000 — which happened after the training traffic — are the untouched test.

4. K-Means clustering (inside each detector)

Addresses → Objective O1 — the specialist-detector design.

Method. Each detector partitions its traffic into clusters so different regions can be judged by different classifiers. K-Means places K centres, assigns each flow to its nearest centre by Euclidean distance ($\sqrt{\Sigma(xi–ci)^2}$), moves each centre to the mean of its flows, and repeats until stable.

In code → src/nids/detector/model.py : DetectorModel.fit() / assign_clusters()

```
self.kmeans = KMeans(self.n_clusters, n_init=10, random_state=seed).fit(Xt_eval)
clusters = self.kmeans.predict(Xt)        # nearest centroid for each flow
```

The code, step by step:

- `n_init=10` restarts K-Means ten times and keeps the best, so one unlucky start doesn't decide the clusters.
- `predict(Xt)` assigns any new flow to the nearest centre — the same operation used at detection time and for the novelty score.

Worked example. With two centres c1 and c2 and a flow at distances 3.2 and 9.8, the flow is put in cluster 1 (3.2 is nearer). After all flows are assigned, c1 moves to the average position of the flows assigned to it.

5. Train the classifiers, select the best per cluster

Addresses → Objective O1 — the specialist-detector design.

Method. Five classifiers (Random Forest, three Decision Trees, Naive Bayes) are trained on the training flows. Then, on the validation flows, each classifier's accuracy is measured per cluster — correct ÷ flows in that cluster — and each cluster keeps the classifier(s) within a small margin α of its best.

In code → src/nids/detector/model.py : fit() + reselect()

```
accuracy = {name: float(np.mean(p[in_cluster] == y_eval[in_cluster]))
            for name, p in predictions.items()}        # accuracy per cluster on VALIDATION
best = max(profile.accuracy.values())
profile.selected = [n for n,a in profile.accuracy.items() if a >= best - alpha]
```

The code, step by step:

- `p[in_cluster] == y_eval[in_cluster]` is a true/false array of whether each validation flow in the cluster was classified correctly; its mean is the accuracy.

- Keeping classifiers within `alpha` of the best lets near-tied classifiers all vote — which is what creates the conflicts that trigger consultation.

Worked example. In one cluster the validation accuracies are RF 0.95, DT1 0.94, DT2 0.90, NB 0.62. With $\alpha = 0.02$ the best is 0.95 and the keep-line is 0.93, so this cluster votes with {RF, DT1} and drops the weaker two.

6. A detector's attack probability and verdict

Addresses → Objective O1 · feeds RQ1 (the ML verdict).

Method. For a flow we find its cluster and average the kept classifiers' attack probabilities, p = (p1+…+pn) ÷ n; the verdict is the majority vote.

In code → src/nids/counselor/decision.py : detector_proba() / _attack_proba() · src/nids/detector/detector.py : detect()

```
probs = np.column_stack([_attack_proba(m.classifiers[n], Xt[rows]) for n in names])
out[rows] = probs.mean(axis=1)            # mean attack probability = p_d(x)
```

The code, step by step:

- `_attack_proba(clf, X)` is a classifier's probability of the **attack** class for every flow in the cluster.
- When a classifier was trained on **benign flows only** — which happens in the leave-one-family-out runs (§11), where a detector's entire attack family is withheld — the attack class is absent from `predict_proba`, so it contributes **0** rather than raising an error. This keeps the unseen-attack experiment runnable.
- `column_stack(...).mean(axis=1)` averages those probabilities — one number per flow.

Worked example. If the kept classifiers output 0.90, 0.70 and 0.80, the detector's probability is (0.90+0.70+0.80)/3 = **0.80**, and the majority vote (all above 0.5) is **attack**.

7. Novelty score and its threshold

Addresses → Objectives O1, O3 · Research question RQ3 (and RQ1) — the unseen-attack signal.

Method. The novelty score is how far a flow sits from the nearest thing the detector learned: n(x) = min over centres of the Euclidean distance. We set τ to the 95th percentile of benign validation scores, so only the 5% most-distant benign flows count as novel. A flow with n(x) > τ is novel — a possible unseen attack.

In code → src/nids/detector/novelty.py : novelty_scores() / choose_threshold()

```
def novelty_scores(detector, X):
    xt = detector.model.transform(X[detector.features])
    return detector.model.kmeans.transform(xt).min(axis=1)   # distance to nearest centre
def choose_threshold(benign_scores, keep_benign=0.95):
    return float(np.quantile(benign_scores, keep_benign))    # tau on benign validation only
```

The code, step by step:

- `kmeans.transform(xt)` returns each flow's distance to every centre; `.min(axis=1)` keeps the nearest — the novelty score.
- `np.quantile(benign_scores, 0.95)` is the value below which 95% of benign scores fall, fixing the benign false-novel rate at ~5%.

Worked example. A flow's distances to three centres are [3.2, 9.8, 7.1], so n(x) = 3.2. If τ = 9.5, then 3.2 < 9.5 — not novel. Another flow at [11.4, 14.0, 12.1] has n(x) = 11.4 > 9.5 — novel.

8. Independent Snort verification

Addresses → Objectives O1, O4 · Research question RQ2 — independent verification.

Method. The ML verdict is stored before Snort is consulted, so Snort never changes it. For each flow we check whether any Snort rule fired, then combine the two independent opinions into Confirmed (both), Disputed (ML normal, Snort fired) or AI-only (ML attack, Snort silent).

In code → scripts/live_detectors/experiment.py : snort_fired() · src/nids/counselor/verification.py

```
snort_fired = df['snort_rules'].astype(str).ne('()')      # did a rule fire on this flow?
def verification_outcome(ml_attack, snort_fired):
    if ml_attack:  return 'confirmed' if snort_fired else 'ai_only'
    return 'disputed' if snort_fired else 'none'
```

The code, step by step:

- `snort_rules` stores the rules that matched; `'()'` means none fired, so `ne('()')` is the Snort verdict.
- The function combines the two verdicts with no feedback between them — this is what makes the verification independent.

Worked example. ML says attack and no rule fired → AI-only (a possible unseen attack). ML says normal but a rule fired → Disputed (Snort caught what the model missed).

9. Alert prioritisation (High / Low)

Addresses → Objective O4 · Research question RQ2 — reducing false positives.

Method. Confirmed, Disputed, an AI-only that is novel or confident (probability ≥ c), and a normal-but-novel "possible unseen attack" are raised as High; weak AI-only and unresolved alerts are logged as Low. The threshold c is set on validation; false positives are measured on High alerts only.

In code → src/nids/counselor/verification.py : priority()

```
if outcome in ('confirmed','disputed'): return 'high'
if outcome == 'ai_only':  return 'high' if (novel or ml_prob >= c) else 'low'
if novel: return 'high'          # ML said normal but the flow is novel
return 'low' if unverified else 'none'
```

The code, step by step:

- Each branch is one row of the design's priority table — the code is the policy.
- Measuring false positives on `'high'` only reflects the alerts an analyst actually acts on.

Worked example. An AI-only flow with probability 0.85 and c = 0.8: 0.85 ≥ 0.8 → High. The same outcome at 0.55 and not novel → Low (logged, not raised to the analyst).

10. Turning counts into answers (metrics)

Addresses → How the numbers for every RQ are computed.

Method. From the four confusion-matrix counts: recall = TP ÷ (TP+FN) is the share of attacks caught; the false-alarm rate = FP ÷ (FP+TN) is the share of benign wrongly raised; MCC combines all four into one score that stays honest under class imbalance.

In code → src/nids/evaluation.py : metrics()

```
recall = tp / (tp + fn)                 # detection rate
false_alarm_rate = fp / (fp + tn)
mcc = matthews_corrcoef(y_true, y_pred) # balances both classes
```

The code, step by step:

- `tp, fp, fn, tn` are counted by comparing the boolean prediction array with the true labels.
- MCC is the headline because accuracy looks high even for a model that just labels everything benign.

Worked example. With TP = 90, FN = 10, FP = 5, TN = 900: recall = 90/(90+10) = 0.90; false-alarm = 5/(5+900) = 0.0055 (5.5 per 1,000); MCC = (90×900 – 5×10) / √(95×100×905×910) ≈ 0.92.

11. Testing RQ1 — leave-one-family-out

Addresses → Objective O3 · Research question RQ1 — unseen attack types.

Method. To test unseen attacks honestly we remove one whole attack family from training and validation, retrain, and test on that family (now unseen) together with benign. Comparing ML-only with the proposed system on that family is the direct answer to RQ1.

In code → scripts/live_detectors/experiment.py : train_detectors() + run()

```
tr2 = tr[~tr.label.isin(withheld)]          # drop the family from training
va2 = va[~va.label.isin(withheld)]          # and from validation
dets = train_detectors(tr2, va2, feats)     # retrain without it
test_unseen = te[te.label.isin(withheld) | (te.label=='Benign')]   # test on it + benign
```

The code, step by step:

- `~...isin(withheld)` keeps every flow whose label is not the withheld family, so the model truly never sees it.
- The test set here is only the withheld family plus benign, so the recall reported is exactly the recall on the unseen family.

Worked example. Withholding DDoS, ML-only caught 27% of the unseen DDoS flows while the full system caught 88% — the gap is the contribution of the independent Snort path and the novelty signal.

12. Testing RQ3 — novelty AUROC

Addresses → Objective O3 · Research question RQ3 — novelty identifies unseen attacks.

Method. AUROC is the probability that a randomly chosen unseen-attack flow has a higher novelty score than a randomly chosen benign flow. 0.5 is a coin flip; higher is better.

In code → src/nids/detector/novelty.py : novelty_auroc()

```
return float(roc_auc_score(is_unseen_attack, scores))   # unseen(1) vs benign(0) by novelty
```

The code, step by step:

- `is_unseen_attack` is the true label (1 unseen attack, 0 benign); `scores` is the novelty score.
- AUROC is threshold-free, so it judges the novelty signal itself, separately from τ.

Worked example. If you draw one unseen-attack flow and one benign flow at random and the attack is more novel 82 times out of 100, the AUROC is 0.82 (our runs gave 0.69–0.93).

13. Testing RQ1 on an unseen network — cross-dataset

Addresses → Objective O3 · Research question RQ1 — an unseen network.

Method. To test generalisation to a network never seen in training, detectors are trained and tuned on one dataset (the source lab) and tested on the other dataset's untouched test partition (the target lab), in both directions: CIC-IDS2017 → CSE-CIC-IDS2018 and CSE-CIC-IDS2018 → CIC-IDS2017. The detectors are fixed — nothing is retrained on the target — so the result measures how well the learned behaviour transfers to a different network. Attack types that appear only in the target (such as HOIC or Zeus) are reported separately, so an unseen network is never confused with an unseen attack type.

In code → scripts/live_detectors/adapt.py : source_detectors() + judge() (the before-adaptation score)

```
dets = source_detectors(flows, source, share=0.6, seed, cache)   # train on the SOURCE lab only
# the target lab's latest 20% per (attack, capture) is held out as a fixed test set
finals, verdict = judge(dets, target_test, snort, trust, min_accuracy)  # score on the unseen network
```

The code, step by step:

- `source_detectors` trains the three detectors on the source lab's earliest 60% and nothing else, so the target network is never in training.
- `judge(...)` runs the whole proposed system — detectors, counsellors and Snort verification — over the target's held-out test flows, and its verdicts are compared with the ML-only and Snort-only baselines on the same flows.

Worked example. Training on CIC-IDS2017 and testing on CSE-CIC-IDS2018, a brute-force flow from the 2018 network — a kind of attack the detector learned on the 2017 network — that is still caught counts towards unseen-network recall; an HOIC flow, which exists only in 2018, is scored separately as an unseen attack type.

14. Testing RQ4 — explanation faithfulness and plausibility

Addresses → Objective O5 · Research question RQ4 — explanation quality.

Method. Faithfulness: we rank a flow's features by their SHAP contribution, neutralise the top-k (set them to their mean), recompute the attack probability and record the drop; we do the same for k random features. If the explanation is faithful, removing the top-k lowers the probability more than removing random features — tested with a Wilcoxon test.
Plausibility: for each attack family a reference list of expected features is fixed in advance, and we report the share of alerts whose five highest-ranked SHAP features include at least one expected feature.

In code → src/nids/detector/faithfulness.py : faithfulness()

```
sv = shap.TreeExplainer(clf).shap_values(x)   # each feature's contribution for each flow
ranked = np.argsort(-np.abs(vals), axis=1)    # rank features by |contribution|
x_top[0, ranked[i,:k]] = 0.0                  # neutralise the top-k
drop = p0[i] - clf.predict_proba(x_top)[0,1]  # fall in attack probability
```

The code, step by step:

- `shap_values` gives, per flow and feature, how much that feature pushed the probability toward attack.
- Setting the chosen features to 0 (their standardised mean) removes their influence without retraining; the drop measures how much they mattered.
- Plausibility checks the ranking against attack knowledge fixed before results are seen — e.g. long flow duration and few bytes per packet for Slowloris — so a faithful-looking ranking must also match how the attack behaves.

Worked example. A flow scores 0.95. Setting its top-5 SHAP features to their mean drops it to 0.39 (a fall of 0.56); setting 5 random features to their mean drops it only to 0.78 (a fall of 0.17). The far larger top-5 fall is the evidence the explanation is faithful.

Counselor NIDS — methodology-and-code reference, aligned with Chapter 3 (four research questions; adaptation is implemented but, as in Chapter 3, left to future work).